

Assignment no 5

Aim:

1. Logistic Regression
2. Differentiate between Linear and Logistic Regression
3. Sigmoid Function
4. Types of LogisticRegression
5. Confusion Matrix Evaluation Metrics

```
In [6]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [7]: df=pd.read_csv("diabetes.csv");
df
```

```
Out[7]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns



```
In [8]: print(df.isnull().sum())
```

```

Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64

```

```

In [6]: cov_matrix = df.cov()
        print(cov_matrix)

```

	Pregnancies	Glucose	BloodPressure	\
Pregnancies	11.354056	13.947131	9.214538	
Glucose	13.947131	1022.248314	94.430956	
BloodPressure	9.214538	94.430956	374.647271	
SkinThickness	-4.390041	29.239183	64.029396	
Insulin	-28.555231	1220.935799	198.378412	
BMI	0.469774	55.726987	43.004695	
DiabetesPedigreeFunction	-0.037426	1.454875	0.264638	
Age	21.570620	99.082805	54.523453	
Outcome	0.356618	7.115079	0.600697	

	SkinThickness	Insulin	BMI	\
Pregnancies	-4.390041	-28.555231	0.469774	
Glucose	29.239183	1220.935799	55.726987	
BloodPressure	64.029396	198.378412	43.004695	
SkinThickness	254.473245	802.979941	49.373869	
Insulin	802.979941	13281.180078	179.775172	
BMI	49.373869	179.775172	62.159984	
DiabetesPedigreeFunction	0.972136	7.066681	0.367405	
Age	-21.381023	-57.143290	3.360330	
Outcome	0.568747	7.175671	1.100638	

	DiabetesPedigreeFunction	Age	Outcome
Pregnancies	-0.037426	21.570620	0.356618
Glucose	1.454875	99.082805	7.115079
BloodPressure	0.264638	54.523453	0.600697
SkinThickness	0.972136	-21.381023	0.568747
Insulin	7.066681	-57.143290	7.175671
BMI	0.367405	3.360330	1.100638
DiabetesPedigreeFunction	0.109779	0.130772	0.027472
Age	0.130772	138.303046	1.336953
Outcome	0.027472	1.336953	0.227483

```

In [9]: X = df.drop('Outcome', axis=1)
        y = df['Outcome']

```

```

In [13]: X

```

Out[13]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 8 columns



In [11]:

y

Out[11]:

```
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
```

Name: Outcome, Length: 768, dtype: int64

In [12]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

In [13]:

```
# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Print the shape of the splits to verify
print(f"Training data shape: {X_train.shape}")
print(f"Testing data shape: {X_test.shape}")
```

Training data shape: (614, 8)

Testing data shape: (154, 8)

```
In [14]: X_train
```

Out[14]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
60	2	84	0	0	0	0.0	
618	9	112	82	24	0	28.2	
346	1	139	46	19	83	28.7	
294	0	161	50	0	0	21.9	
231	6	134	80	37	370	46.2	
...	
71	5	139	64	35	140	28.6	
106	1	96	122	0	0	22.4	
270	10	101	86	37	0	45.6	
435	0	141	0	0	0	42.4	
102	0	125	96	0	0	22.5	

614 rows × 8 columns



```
In [15]: X_test
```

Out[15]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
668	6	98	58	33	190	34.0	
324	2	112	75	32	0	35.7	
624	2	108	64	0	0	30.8	
690	8	107	80	0	0	24.6	
473	7	136	90	0	0	29.9	
...	
355	9	165	88	0	0	30.4	
534	1	77	56	30	56	33.3	
344	8	95	72	0	0	36.8	
296	2	146	70	38	360	28.0	
462	8	74	70	40	49	35.3	

154 rows × 8 columns



```
In [16]: y_train
```

```
Out[16]: 60      0
        618     1
        346     0
        294     0
        231     1
          ..
        71      0
       106      0
       270      1
       435      1
       102      0
        Name: Outcome, Length: 614, dtype: int64
```

```
In [17]: y_test
```

```
Out[17]: 668      0
        324      0
        624      0
        690      0
        473      0
          ..
       355      1
       534      0
       344      0
       296      1
       462      0
        Name: Outcome, Length: 154, dtype: int64
```

```
In [18]: # Initialize the StandardScaler
        scaler = StandardScaler()

        # Fit on training data and transform both training and testing data
        X_train = scaler.fit_transform(X_train)
        X_test = scaler.transform(X_test)
```

```
In [19]: X_train
```

```
Out[19]: array([[ -0.52639686, -1.15139792, -3.75268255, ..., -4.13525578,
        -0.49073479, -1.03594038],
        [ 1.58804586, -0.27664283,  0.68034485, ..., -0.48916881,
        2.41502991,  1.48710085],
        [-0.82846011,  0.56687102, -1.2658623 , ..., -0.42452187,
        0.54916055, -0.94893896],
        ...,
        [ 1.8901091 , -0.62029661,  0.89659009, ...,  1.76054443,
        1.981245  ,  0.44308379],
        [-1.13052335,  0.62935353, -3.75268255, ...,  1.34680407,
        -0.78487662, -0.33992901],
        [-1.13052335,  0.12949347,  1.43720319, ..., -1.22614383,
        -0.61552223, -1.03594038]])
```

```
In [20]: X_test
```

```
Out[20]: array([[ 0.68185612, -0.71402038, -0.61712658, ...,  0.26073561,
                 -0.11637247,  0.87809089],
                [-0.52639686, -0.27664283,  0.30191569, ...,  0.48053518,
                 -0.954231  , -1.03594038],
                [-0.52639686, -0.40160784, -0.29275872, ..., -0.15300476,
                 -0.9245197  , -1.03594038],
                ...,
                [ 1.28598261, -0.80774414,  0.13973176, ...,  0.62275843,
                 0.04703966,  2.0961108  ],
                [-0.52639686,  0.78555979,  0.03160914, ..., -0.51502758,
                 -0.39268751, -0.33992901],
                [ 1.28598261, -1.46381046,  0.03160914, ...,  0.42881763,
                 0.70068816,  0.53008521]])
```

```
In [21]: # Initialize the Logistic Regression model
logreg = LogisticRegression(max_iter=1000)

# Fit the model to the training data
logreg.fit(X_train, y_train)
```

```
Out[21]: ▼      LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [24]: # Predict on the training data
y_train_pred = logreg.predict(X_train)

# Predict on the testing data
y_test_pred = logreg.predict(X_test)
```

```
In [25]: y_train_pred
```

```
Out[25]: array([0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0,
               0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0,
               0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
               1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
               1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0,
               1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1,
               0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0,
               1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
               1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1,
               1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
               1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
               0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1,
               0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0,
               1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0,
               0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0,
               0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
               0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
               0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
               1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
               0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
               0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
               0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
               0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0],
               dtype=int64)
```

```
In [27]: y_test_pred
```

```
Out[27]: array([0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0,
               1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
               0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1,
               0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0,
               0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1,
               0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1,
               0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0],
               dtype=int64)
```

```
In [29]: train_accuracy = accuracy_score(y_train, y_train_pred)
         train_accuracy
```

```
Out[29]: 0.7703583061889251
```

```
In [30]: train_precision = precision_score(y_train, y_train_pred)
         train_precision
```

```
Out[30]: 0.7168674698795181
```

```
In [31]: train_recall = recall_score(y_train, y_train_pred)
         train_recall
```

```
Out[31]: 0.5586854460093896
```

```
In [32]: train_f1 = f1_score(y_train, y_train_pred)
train_f1
```

```
Out[32]: 0.6279683377308708
```

```
In [33]: train_cm = confusion_matrix(y_train, y_train_pred)
train_cm
```

```
Out[33]: array([[354,  47],
               [ 94, 119]], dtype=int64)
```

```
In [34]: test_accuracy = accuracy_score(y_test, y_test_pred)
test_accuracy
```

```
Out[34]: 0.7532467532467533
```

```
In [35]: test_precision = precision_score(y_test, y_test_pred)
test_precision
```

```
Out[35]: 0.6491228070175439
```

```
In [36]: test_recall = recall_score(y_test, y_test_pred)
test_recall
```

```
Out[36]: 0.6727272727272727
```

```
In [37]: test_f1 = f1_score(y_test, y_test_pred)
test_f1
```

```
Out[37]: 0.6607142857142858
```

```
In [38]: test_cm = confusion_matrix(y_test, y_test_pred)
test_cm
```

```
Out[38]: array([[79, 20],
               [18, 37]], dtype=int64)
```

```
In [ ]: Name:Kadhane Pratiksha
Rollno:13213
Batch:B1
```